

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

# Cleaned data from the wellbeing and lifestyle kaggle dataset
data = pd.read_csv("data_clean.csv") # Unscaled
# Cleaned training, validation, and test data from same dataset
train_data = pd.read_csv("data_train.csv")
val_data = pd.read_csv("data_val.csv")
test_data = pd.read_csv("data_test.csv")

train_data.head()

```

	FRUITS_VEGGIES	DAILY_STRESS	PLACES_VISITED	CORE_CIRCLE	\
0	-0.636358	0.867892	-0.376228	0.519366	
1	-1.328086	0.139683	1.431552	1.569497	
2	-0.636358	0.867892	-0.376228	-0.530764	
3	0.747098	-1.316734	-0.677524	0.869410	
4	0.747098	-1.316734	1.431552	1.569497	

	SUPPORTING_OTHERS	SOCIAL_NETWORK	ACHIEVEMENT	DONATION	BMI_RANGE	\
0	0.416268	1.138570	0.002509	0.688858	-0.836204	
1	-0.815281	-0.157744	-0.726154	1.230072	-0.836204	
2	-0.507394	-0.157744	0.731173	0.688858	1.195880	
3	0.724156	1.138570	1.824168	1.230072	-0.836204	
4	1.339930	0.490413	1.095505	1.230072	-0.836204	

	TODO_COMPLETED	...	PERSONAL_AWARDS	TIME_FOR_PASSION	WEEKLY_MEDITATION	\
0	0.097801	...	-0.231363	-0.476632	-0.738927	
1	0.097801	...	-0.231363	-1.210773	0.253374	
2	0.097801	...	1.392280	-0.843703	-0.738927	
3	-1.042719	...	1.392280	0.624579	1.245675	
4	0.477974	...	-0.880820	0.991650	0.253374	

	WORK_LIFE_BALANCE_SCORE	AGE_21 to 35	AGE_36 to 50	AGE_51 or more	\
0	0.463390	0.0	0.0	0.0	
1	0.117818	0.0	1.0	0.0	
2	-0.455164	0.0	0.0	1.0	
3	1.121093	0.0	0.0	1.0	
4	2.115450	0.0	0.0	1.0	

	AGE_Less than 20	GENDER_Female	GENDER_Male
0	1.0	1.0	0.0
1	0.0	1.0	0.0
2	0.0	1.0	0.0
3	0.0	1.0	0.0
4	0.0	1.0	0.0

[5 rows x 27 columns]

val_data.head()

	FRUITS_VEGGIES	DAILY_STRESS	PLACES_VISITED	CORE_CIRCLE	\
0	0.055370	0.139683	-1.581414	-1.230851	
1	0.747098	-0.588525	-0.677524	1.219453	
2	0.055370	0.139683	0.226365	-0.180721	
3	0.055370	-1.316734	-0.074931	0.519366	
4	0.055370	0.867892	-1.280117	1.569497	

	SUPPORTING_OTHERS	SOCIAL_NETWORK	ACHIEVEMENT	DONATION	BMI_RANGE	\
0	-1.123169	-1.454058	-1.454817	-1.475996	-0.836204	
1	1.339930	1.138570	1.459836	-0.393569	-0.836204	
2	1.339930	1.138570	0.731173	1.230072	-0.836204	
3	-0.507394	-1.129979	-0.726154	-1.475996	-0.836204	
4	1.032043	0.490413	0.366841	-0.934782	1.195880	

	TODO_COMPLETED	...	PERSONAL_AWARDS	TIME_FOR_PASSION	WEEKLY_MEDITATION	\
0	0.858147	...	-1.205549	-1.210773	0.253374	
1	0.858147	...	0.418094	0.257509	1.245675	
2	-0.282372	...	-0.880820	0.991650	1.245675	
3	0.097801	...	-0.880820	-1.210773	1.245675	
4	0.477974	...	-0.880820	-0.843703	1.245675	

	WORK_LIFE_BALANCE_SCORE	AGE_21 to 35	AGE_36 to 50	AGE_51 or more	\
0	-1.048211	0.0	0.0	1.0	
1	1.103257	0.0	0.0	1.0	
2	1.315060	0.0	1.0	0.0	
3	0.371981	0.0	0.0	0.0	
4	-0.660278	0.0	1.0	0.0	

	AGE_Less than 20	GENDER_Female	GENDER_Male
0	0.0	1.0	0.0
1	0.0	0.0	1.0
2	0.0	0.0	1.0
3	1.0	0.0	1.0
4	0.0	1.0	0.0

[5 rows x 27 columns]

test_data.head()

	FRUITS_VEGGIES	DAILY_STRESS	PLACES_VISITED	CORE_CIRCLE	\
0	-0.636358	0.139683	-0.376228	-0.880808	
1	0.055370	-0.588525	-0.074931	0.169323	
2	-0.636358	0.139683	-1.581414	-1.230851	

3	-1.328086	0.867892	-0.677524	-1.580895
4	-1.328086	1.596101	-0.978821	-1.580895

	SUPPORTING_OTHERS	SOCIAL_NETWORK	ACHIEVEMENT	DONATION	BMI_RANGE \
0	-1.431056	-0.805901	-1.090486	-1.475996	-0.836204
1	-1.738944	1.138570	-0.361822	-0.393569	-0.836204
2	1.339930	-1.454058	1.095505	1.230072	1.195880
3	-0.815281	-0.481822	-0.361822	-0.934782	-0.836204
4	-1.738944	-1.778136	-1.090486	-1.475996	-0.836204

	TODO_COMPLETED ...	PERSONAL_AWARDS	TIME_FOR_PASSION	WEEKLY_MEDITATION \
0	0.097801 ...	-0.880820	-0.109562	-0.408160
1	-1.422892 ...	-0.231363	-0.843703	0.584141
2	0.477974 ...	0.093365	-0.476632	-1.400461
3	-1.422892 ...	-0.880820	-0.843703	-1.400461
4	-0.662546 ...	-0.880820	-1.210773	-1.731228

	WORK_LIFE_BALANCE_SCORE	AGE_21 to 35	AGE_36 to 50	AGE_51 or more \
0	-1.019228	0.0	0.0	0.0
1	-0.716016	1.0	0.0	0.0
2	-0.247821	0.0	0.0	1.0
3	-1.625652	1.0	0.0	0.0
4	-2.468403	1.0	0.0	0.0

	AGE_Less than 20	GENDER_Female	GENDER_Male
0	1.0	0.0	1.0
1	0.0	1.0	0.0
2	0.0	0.0	1.0
3	0.0	1.0	0.0
4	0.0	1.0	0.0

[5 rows x 27 columns]

Task 2: Build and Train Various ML Models

Part 1: Logistic Regression

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
```

Prepare the datasets

```
X_train = train_data.drop(['GENDER_Female', 'GENDER_Male'], axis=1)
y_train = train_data['GENDER_Male']
X_val = val_data.drop(['GENDER_Female', 'GENDER_Male'], axis=1)
y_val = val_data['GENDER_Male']
X_test = test_data.drop(['GENDER_Female', 'GENDER_Male'], axis=1)
```

```

y_test = test_data['GENDER_Male']

# Initialize and train the logistic regression model
max_iter_val = 5000
# Adjusting the hyperparameters to what gave the highest found accuracy (based on cell below)
log_reg = LogisticRegression(max_iter=max_iter_val)
log_reg.fit(X_train, y_train)

y_train_pred = log_reg.predict(X_train)
train_accuracy = accuracy_score(y_train, y_train_pred)
print(f"Training Accuracy: {train_accuracy}")

# Predict on the validation set
y_val_pred = log_reg.predict(X_val)

# Evaluate the model on the validation set
val_accuracy = accuracy_score(y_val, y_val_pred)
val_conf_matrix = confusion_matrix(y_val, y_val_pred)

print(f"Validation Accuracy: {val_accuracy}")
print("Confusion Matrix:\n", val_conf_matrix)

Training Accuracy: 0.6669797537048633
Validation Accuracy: 0.6651017214397497
Confusion Matrix:
[[1701  264]
 [ 806  424]]

# Trying out different hyperparameters
C_values = [0.001, 0.01, 0.1, 1, 10, 100]
penalties = ['l1', 'l2', 'elasticnet', 'none']

from sklearn.metrics import accuracy_score

best_accuracy = 0
best_params = {'C': None, 'penalty': None}

for penalty in penalties:
    # Skip invalid solver/penalty combinations
    if penalty == 'elasticnet':
        solver = 'saga'
    else:
        solver = 'liblinear'

    for C in C_values:
        # Initialize the Logistic Regression model with current parameters
        if penalty == 'elasticnet':

```

```

        model = LogisticRegression(C=C, penalty=penalty, l1_ratio=0.5, solver=solver, max_iter=1000)
    else:
        model = LogisticRegression(C=C, penalty=penalty, solver=solver, max_iter=1000)

    # Fitting the model on training data and predicting on validation data
    try:
        model.fit(X_train, y_train)
        y_val_pred = model.predict(X_val)
        accuracy = accuracy_score(y_val, y_val_pred)

        # printing performance and update the best parameters
        print(f"Penalty: {penalty}, C: {C}, Accuracy: {accuracy}")
        if accuracy > best_accuracy:
            best_accuracy = accuracy
            best_params['C'] = C
            best_params['penalty'] = penalty
    except Exception as e:
        # This block catches errors which might arise from invalid parameter combination
        print(f"Could not train model with Penalty: {penalty}, C: {C}. Error: {str(e)}")

best_model = LogisticRegression(C=best_params['C'], penalty=best_params['penalty'], solver=
best_model.fit(X_train, y_train)
y_train_pred = best_model.predict(X_train)
train_accuracy = accuracy_score(y_train, y_train_pred)
print(f"SVM Training Accuracy: {train_accuracy}")
# Print out the best parameters and best accuracy found
print(f"Best parameters: {best_params}")
print(f"Best validation accuracy: {best_accuracy}")

Penalty: 11, C: 0.001, Accuracy: 0.6150234741784038
Penalty: 11, C: 0.01, Accuracy: 0.6613458528951487
Penalty: 11, C: 0.1, Accuracy: 0.6613458528951487
Penalty: 11, C: 1, Accuracy: 0.6644757433489827
Penalty: 11, C: 10, Accuracy: 0.6651017214397497
Penalty: 11, C: 100, Accuracy: 0.6651017214397497
Penalty: 12, C: 0.001, Accuracy: 0.6629107981220658
Penalty: 12, C: 0.01, Accuracy: 0.6632237871674491
Penalty: 12, C: 0.1, Accuracy: 0.6644757433489827
Penalty: 12, C: 1, Accuracy: 0.6651017214397497
Penalty: 12, C: 10, Accuracy: 0.6651017214397497
Penalty: 12, C: 100, Accuracy: 0.6651017214397497
Penalty: elasticnet, C: 0.001, Accuracy: 0.6150234741784038
Penalty: elasticnet, C: 0.01, Accuracy: 0.662284820031299
Penalty: elasticnet, C: 0.1, Accuracy: 0.6632237871674491
Penalty: elasticnet, C: 1, Accuracy: 0.6647887323943662
Penalty: elasticnet, C: 10, Accuracy: 0.6651017214397497

```

```

/Users/madmax11/anaconda3/lib/python3.11/site-packages/sklearn/linear_model/_sag.py:350: Con
warnings.warn(

Penalty: elasticnet, C: 100, Accuracy: 0.6651017214397497
Could not train model with Penalty: none, C: 0.001. Error: penalty='none' is not supported f
Could not train model with Penalty: none, C: 0.01. Error: penalty='none' is not supported fo
Could not train model with Penalty: none, C: 0.1. Error: penalty='none' is not supported for
Could not train model with Penalty: none, C: 1. Error: penalty='none' is not supported for t
Could not train model with Penalty: none, C: 10. Error: penalty='none' is not supported for
Could not train model with Penalty: none, C: 100. Error: penalty='none' is not supported for
SVM Training Accuracy: 0.6669797537048633
Best parameters: {'C': 10, 'penalty': 'l1'}
Best validation accuracy: 0.6651017214397497

#Optimized model with new beset found hyper parameters
log_reg_optimized = LogisticRegression(max_iter=5000, C=0.01, penalty='elasticnet', solver=
log_reg_optimized.fit(X_train, y_train)

#Prediction on test set
y_test_pred = log_reg_optimized.predict(X_test)

# Evaluate the model on the validation set
test_accuracy = accuracy_score(y_val, y_val_pred)
test_conf_matrix = confusion_matrix(y_val, y_val_pred)


print(f"Test Accuracy: {test_accuracy}")
print("Confusion Matrix:\n", test_conf_matrix)

plt.figure(figsize=(7, 5))
sns.heatmap(val_conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['Female', 'Male
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix for Test Data')

Test Accuracy: 0.6651017214397497
Confusion Matrix:
[[1701  264]
 [ 806  424]]

Text(0.5, 1.0, 'Confusion Matrix for Test Data')

```



1. Support Vector Machine

```

from sklearn.svm import SVC

# Training model and checking different hyperparameters
kernels = ['linear', 'rbf', 'poly']
for kernel in kernels:
    svm_model = SVC(kernel=kernel, C=1.0)
    if kernel == 'poly':
        svm_model.degree = 3
    svm_model.fit(X_train, y_train)
    y_val_pred = svm_model.predict(X_val)
    val_accuracy = accuracy_score(y_val, y_val_pred)
    print(f'Validation Accuracy with {kernel} kernel: {val_accuracy}')

Validation Accuracy with linear kernel: 0.6528951486697966
Validation Accuracy with rbf kernel: 0.6666666666666666
Validation Accuracy with poly kernel: 0.6519561815336463

#Optimized model
svm_optimized_model = SVC(kernel='rbf', C=1.0, gamma='scale')
svm_optimized_model.fit(X_train, y_train)

y_test_pred = svm_optimized_model.predict(X_test)

```

```

test_accuracy = accuracy_score(y_test, y_test_pred)
test_conf_matrix = confusion_matrix(y_test, y_test_pred)


print(f"Test Accuracy: {test_accuracy}")
print("Confusion Matrix for test data:\n", test_conf_matrix)

plt.figure(figsize=(7, 5))
sns.heatmap(test_conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['Female', 'Male'], yticklabels=['Female', 'Male'])
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix for SVM on Test Data')

Test Accuracy: 0.6685446009389672
Confusion Matrix for test data:
[[1737  198]
 [ 861  399]]

Text(0.5, 1.0, 'Confusion Matrix for SVM on Test Data')

```



1. Decision Trees

```

age_columns = ['AGE_21 to 35', 'AGE_36 to 50', 'AGE_51 or more', 'AGE_Less than 20']

for df in [train_data, val_data, test_data]:

```



```

conditions = [
    df['AGE_21 to 35'] == 1,
    df['AGE_36 to 50'] == 1,
    df['AGE_51 or more'] == 1,
    df['AGE_Less than 20'] == 1
]
# The choices must correspond one-to-one with the conditions
choices = ['21-35', '36-50', '51+', '<20']
df['age_group'] = np.select(conditions, choices, default='Unknown')

# Now, prepare your feature and target datasets
X_train = train_data.drop(age_columns + ['age_group'], axis=1)
y_train = train_data['age_group']
X_val = val_data.drop(age_columns + ['age_group'], axis=1)
y_val = val_data['age_group']
X_test = test_data.drop(age_columns + ['age_group'], axis=1)
y_test = test_data['age_group']

from sklearn.tree import DecisionTreeClassifier

dt_classifier = DecisionTreeClassifier(random_state=11)
dt_classifier.fit(X_train, y_train)

DecisionTreeClassifier(random_state=11)

from sklearn.tree import plot_tree

y_val_pred = dt_classifier.predict(X_val)

y_train_pred = dt_classifier.predict(X_train)
train_accuracy = accuracy_score(y_train, y_train_pred)

val_accuracy = accuracy_score(y_val, y_val_pred)
val_conf_matrix = confusion_matrix(y_val, y_val_pred)
val_class_report = classification_report(y_val, y_val_pred)

print(f"Training Accuracy: {train_accuracy}")
print(f"Validation Accuracy: {val_accuracy}")
print("Confusion Matrix:\n", val_conf_matrix)
print("Classification Report:\n", val_class_report)

# This takes about 2 minutes
# plt.figure(figsize=(20,10))
# plot_tree(dt_classifier, feature_names=X_train.columns, class_names=age_groups, filled=True)

Training Accuracy: 0.9996869129618033
Validation Accuracy: 0.38059467918622847
Confusion Matrix:

```

```

[[557 340 200 142]
 [287 329 194 88]
 [202 205 241 53]
 [127 82 59 89]]
Classification Report:
              precision    recall  f1-score   support

   21-35         0.47         0.45         0.46         1239
   36-50         0.34         0.37         0.35          898
    51+         0.35         0.34         0.35          701
    <20         0.24         0.25         0.24          357

 accuracy                   0.38         3195
 macro avg              0.35         0.35         0.35         3195
weighted avg              0.38         0.38         0.38         3195

y_test_pred = dt_classifier.predict(X_test)

test_accuracy = accuracy_score(y_test, y_test_pred)
test_conf_matrix = confusion_matrix(y_test, y_test_pred)

print(f"Test Accuracy: {test_accuracy}")
print("Confusion Matrix on Test Data:\n", test_conf_matrix)

Test Accuracy: 0.38841940532081376
Confusion Matrix on Test Data:
[[552 307 202 138]
 [301 356 197 85]
 [191 195 232 60]
 [132 86 60 101]]

```

1. Random Forests

```

from sklearn.ensemble import RandomForestClassifier

rf_classifier = RandomForestClassifier(random_state=17)
rf_classifier.fit(X_train, y_train)

RandomForestClassifier(random_state=17)

y_val_pred = rf_classifier.predict(X_val)
y_train_pred = rf_classifier.predict(X_train)

# Evaluating the model
train_accuracy = accuracy_score(y_train, y_train_pred)
val_accuracy = accuracy_score(y_val, y_val_pred)
val_conf_matrix = confusion_matrix(y_val, y_val_pred)
val_class_report = classification_report(y_val, y_val_pred)

```

```

print(f"Random Forest Training Accuracy: {train_accuracy}")
print(f"Validation Accuracy: {val_accuracy}")
print("Confusion Matrix:\n", val_conf_matrix)
print("Classification Report:\n", val_class_report)

```

Random Forest Training Accuracy: 0.9996869129618033

Validation Accuracy: 0.48012519561815337

Confusion Matrix:

```

[[914 228 82 15]
 [433 338 117 10]
 [271 187 241 2]
 [243 58 15 41]]

```

Classification Report:

	precision	recall	f1-score	support
21-35	0.49	0.74	0.59	1239
36-50	0.42	0.38	0.40	898
51+	0.53	0.34	0.42	701
<20	0.60	0.11	0.19	357
accuracy			0.48	3195
macro avg	0.51	0.39	0.40	3195
weighted avg	0.49	0.48	0.45	3195

```

y_test_pred = rf_classifier.predict(X_test)

```

Evaluating the model on the test set

```

test_accuracy = accuracy_score(y_test, y_test_pred)

```

```

test_conf_matrix = confusion_matrix(y_test, y_test_pred)

```

```

print(f"Test Accuracy: {test_accuracy}")

```

```

print("Confusion Matrix on Test Data:\n", test_conf_matrix)

```

Test Accuracy: 0.48200312989045385

Confusion Matrix on Test Data:

```

[[897 208 74 20]
 [458 356 116 9]
 [258 192 227 1]
 [248 53 18 60]]

```

```

plt.figure(figsize=(20,10))

```

```

plot_tree(rf_classifier.estimators_[0], feature_names=X_train.columns, class_names=age_column_names)

```

[Text(0.5, 0.9, 'BMI_RANGE <= 0.18\ngini = 0.707\nsamples = 6021\nvalue = [3711, 2844, 1970, 1414])

Text(0.25, 0.7, 'WORK_LIFE_BALANCE_SCORE <= 0.293\ngini = 0.7\nsamples = 3502\nvalue = [242, 108, 1970, 1414])

Text(0.125, 0.5, 'SOCIAL_NETWORK <= 0.004\ngini = 0.671\nsamples = 1884\nvalue = [1495, 709, 1970, 1414])


```

# Reload the data
train_data = pd.read_csv("data_train.csv")
val_data = pd.read_csv("data_val.csv")
test_data = pd.read_csv("data_test.csv")

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
import tensorflow as tf
from sklearn.metrics import mean_squared_error

X_train = train_data.drop('WORK_LIFE_BALANCE_SCORE', axis=1)
y_train = train_data['WORK_LIFE_BALANCE_SCORE']
X_val = val_data.drop('WORK_LIFE_BALANCE_SCORE', axis=1)
y_val = val_data['WORK_LIFE_BALANCE_SCORE']
X_test = test_data.drop('WORK_LIFE_BALANCE_SCORE', axis=1)
y_test = test_data['WORK_LIFE_BALANCE_SCORE']

def build_model(input_shape):
    # Initialize the neural net structure
    model = Sequential([
        Dense(64, activation='relu', input_shape=(input_shape,)),
        Dense(64, activation='relu'),
        Dense(1)
    ])
    # Compile the model
    model.compile(optimizer='adam', loss='mse')
    return model

num_seeds = 10
all_scoresT = []
all_scoresV = []

# n different seeds for training
for i in range(num_seeds):
    print(f"Training on seed {i}")
    tf.random.set_seed(i)
    model = build_model(X_train.shape[1])
    model.fit(X_train, y_train, epochs=10, batch_size=32, verbose=1)

    y_train_pred = model.predict(X_train)
    train_score = np.sqrt(mean_squared_error(y_train, y_train_pred))
    all_scoresT.append(train_score)

    y_val_pred = model.predict(X_val)
    val_score = np.sqrt(mean_squared_error(y_val, y_val_pred))
    all_scoresV.append(val_score)

```

```

average_scoreT = np.mean(all_scoresT)
average_scoreV = np.mean(all_scoresV)
print(f"Average Training Score: {average_scoreT}")
print(f"Average Validation Score: {average_scoreV}")

Training on seed 0
Epoch 1/10
300/300 [=====] - 0s 538us/step - loss: 0.0573
Epoch 2/10
300/300 [=====] - 0s 513us/step - loss: 0.0075
Epoch 3/10
300/300 [=====] - 0s 517us/step - loss: 0.0042
Epoch 4/10
300/300 [=====] - 0s 533us/step - loss: 0.0028
Epoch 5/10
300/300 [=====] - 0s 540us/step - loss: 0.0020
Epoch 6/10
300/300 [=====] - 0s 534us/step - loss: 0.0015
Epoch 7/10
300/300 [=====] - 0s 535us/step - loss: 0.0012
Epoch 8/10
300/300 [=====] - 0s 524us/step - loss: 9.6339e-04
Epoch 9/10
300/300 [=====] - 0s 539us/step - loss: 8.2076e-04
Epoch 10/10
300/300 [=====] - 0s 518us/step - loss: 6.6603e-04
300/300 [=====] - 0s 336us/step
100/100 [=====] - 0s 342us/step
Training on seed 1
Epoch 1/10
300/300 [=====] - 0s 534us/step - loss: 0.0531
Epoch 2/10
300/300 [=====] - 0s 562us/step - loss: 0.0074
Epoch 3/10
300/300 [=====] - 0s 557us/step - loss: 0.0041
Epoch 4/10
300/300 [=====] - 0s 557us/step - loss: 0.0028
Epoch 5/10
300/300 [=====] - 0s 542us/step - loss: 0.0020
Epoch 6/10
300/300 [=====] - 0s 539us/step - loss: 0.0016
Epoch 7/10
300/300 [=====] - 0s 537us/step - loss: 0.0013
Epoch 8/10
300/300 [=====] - 0s 539us/step - loss: 0.0010
Epoch 9/10

```

```

300/300 [=====] - 0s 538us/step - loss: 8.5368e-04
Epoch 10/10
300/300 [=====] - 0s 531us/step - loss: 6.9508e-04
300/300 [=====] - 0s 349us/step
100/100 [=====] - 0s 365us/step
Training on seed 2
Epoch 1/10
300/300 [=====] - 0s 551us/step - loss: 0.0583
Epoch 2/10
300/300 [=====] - 0s 540us/step - loss: 0.0069
Epoch 3/10
300/300 [=====] - 0s 532us/step - loss: 0.0038
Epoch 4/10
300/300 [=====] - 0s 535us/step - loss: 0.0025
Epoch 5/10
300/300 [=====] - 0s 536us/step - loss: 0.0018
Epoch 6/10
300/300 [=====] - 0s 531us/step - loss: 0.0014
Epoch 7/10
300/300 [=====] - 0s 541us/step - loss: 0.0011
Epoch 8/10
300/300 [=====] - 0s 535us/step - loss: 8.1736e-04
Epoch 9/10
300/300 [=====] - 0s 524us/step - loss: 6.4278e-04
Epoch 10/10
300/300 [=====] - 0s 520us/step - loss: 5.2958e-04
300/300 [=====] - 0s 351us/step
100/100 [=====] - 0s 367us/step
Training on seed 3
Epoch 1/10
300/300 [=====] - 0s 555us/step - loss: 0.0562
Epoch 2/10
300/300 [=====] - 0s 545us/step - loss: 0.0079
Epoch 3/10
300/300 [=====] - 0s 543us/step - loss: 0.0047
Epoch 4/10
300/300 [=====] - 0s 544us/step - loss: 0.0031
Epoch 5/10
300/300 [=====] - 0s 539us/step - loss: 0.0024
Epoch 6/10
300/300 [=====] - 0s 529us/step - loss: 0.0017
Epoch 7/10
300/300 [=====] - 0s 531us/step - loss: 0.0014
Epoch 8/10
300/300 [=====] - 0s 546us/step - loss: 0.0011
Epoch 9/10

```

```

300/300 [=====] - 0s 546us/step - loss: 8.9771e-04
Epoch 10/10
300/300 [=====] - 0s 544us/step - loss: 7.5501e-04
300/300 [=====] - 0s 356us/step
100/100 [=====] - 0s 356us/step
Training on seed 4
Epoch 1/10
300/300 [=====] - 0s 560us/step - loss: 0.0668
Epoch 2/10
300/300 [=====] - 0s 538us/step - loss: 0.0066
Epoch 3/10
300/300 [=====] - 0s 536us/step - loss: 0.0037
Epoch 4/10
300/300 [=====] - 0s 556us/step - loss: 0.0025
Epoch 5/10
300/300 [=====] - 0s 528us/step - loss: 0.0018
Epoch 6/10
300/300 [=====] - 0s 528us/step - loss: 0.0014
Epoch 7/10
300/300 [=====] - 0s 541us/step - loss: 0.0010
Epoch 8/10
300/300 [=====] - 0s 551us/step - loss: 8.5196e-04
Epoch 9/10
300/300 [=====] - 0s 549us/step - loss: 7.2771e-04
Epoch 10/10
300/300 [=====] - 0s 542us/step - loss: 5.6479e-04
300/300 [=====] - 0s 358us/step
100/100 [=====] - 0s 363us/step
Training on seed 5
Epoch 1/10
300/300 [=====] - 0s 562us/step - loss: 0.0693
Epoch 2/10
300/300 [=====] - 0s 594us/step - loss: 0.0080
Epoch 3/10
300/300 [=====] - 0s 537us/step - loss: 0.0044
Epoch 4/10
300/300 [=====] - 0s 533us/step - loss: 0.0030
Epoch 5/10
300/300 [=====] - 0s 532us/step - loss: 0.0022
Epoch 6/10
300/300 [=====] - 0s 538us/step - loss: 0.0017
Epoch 7/10
300/300 [=====] - 0s 533us/step - loss: 0.0013
Epoch 8/10
300/300 [=====] - 0s 533us/step - loss: 0.0011
Epoch 9/10

```



```

300/300 [=====] - 0s 539us/step - loss: 8.4299e-04
Epoch 10/10
300/300 [=====] - 0s 540us/step - loss: 6.9731e-04
300/300 [=====] - 0s 359us/step
100/100 [=====] - 0s 350us/step
Training on seed 6
Epoch 1/10
300/300 [=====] - 0s 551us/step - loss: 0.0636
Epoch 2/10
300/300 [=====] - 0s 551us/step - loss: 0.0086
Epoch 3/10
300/300 [=====] - 0s 535us/step - loss: 0.0048
Epoch 4/10
300/300 [=====] - 0s 538us/step - loss: 0.0032
Epoch 5/10
300/300 [=====] - 0s 535us/step - loss: 0.0023
Epoch 6/10
300/300 [=====] - 0s 535us/step - loss: 0.0017
Epoch 7/10
300/300 [=====] - 0s 535us/step - loss: 0.0014
Epoch 8/10
300/300 [=====] - 0s 531us/step - loss: 0.0011
Epoch 9/10
300/300 [=====] - 0s 533us/step - loss: 8.6700e-04
Epoch 10/10
300/300 [=====] - 0s 560us/step - loss: 7.0429e-04
300/300 [=====] - 0s 357us/step
100/100 [=====] - 0s 366us/step
Training on seed 7
Epoch 1/10
300/300 [=====] - 0s 544us/step - loss: 0.0741
Epoch 2/10
300/300 [=====] - 0s 548us/step - loss: 0.0086
Epoch 3/10
300/300 [=====] - 0s 531us/step - loss: 0.0051
Epoch 4/10
300/300 [=====] - 0s 541us/step - loss: 0.0035
Epoch 5/10
300/300 [=====] - 0s 561us/step - loss: 0.0025
Epoch 6/10
300/300 [=====] - 0s 546us/step - loss: 0.0020
Epoch 7/10
300/300 [=====] - 0s 540us/step - loss: 0.0016
Epoch 8/10
300/300 [=====] - 0s 539us/step - loss: 0.0013
Epoch 9/10

```

```

300/300 [=====] - 0s 537us/step - loss: 0.0010
Epoch 10/10
300/300 [=====] - 0s 545us/step - loss: 9.0563e-04
300/300 [=====] - 0s 346us/step
100/100 [=====] - 0s 368us/step
Training on seed 8
Epoch 1/10
300/300 [=====] - 0s 573us/step - loss: 0.0721
Epoch 2/10
300/300 [=====] - 0s 537us/step - loss: 0.0088
Epoch 3/10
300/300 [=====] - 0s 533us/step - loss: 0.0051
Epoch 4/10
300/300 [=====] - 0s 526us/step - loss: 0.0034
Epoch 5/10
300/300 [=====] - 0s 530us/step - loss: 0.0024
Epoch 6/10
300/300 [=====] - 0s 527us/step - loss: 0.0019
Epoch 7/10
300/300 [=====] - 0s 527us/step - loss: 0.0015
Epoch 8/10
300/300 [=====] - 0s 532us/step - loss: 0.0012
Epoch 9/10
300/300 [=====] - 0s 530us/step - loss: 0.0010
Epoch 10/10
300/300 [=====] - 0s 529us/step - loss: 8.3641e-04
300/300 [=====] - 0s 342us/step
100/100 [=====] - 0s 352us/step
Training on seed 9
Epoch 1/10
300/300 [=====] - 0s 577us/step - loss: 0.0569
Epoch 2/10
300/300 [=====] - 0s 570us/step - loss: 0.0079
Epoch 3/10
300/300 [=====] - 0s 545us/step - loss: 0.0045
Epoch 4/10
300/300 [=====] - 0s 571us/step - loss: 0.0029
Epoch 5/10
300/300 [=====] - 0s 536us/step - loss: 0.0020
Epoch 6/10
300/300 [=====] - 0s 536us/step - loss: 0.0015
Epoch 7/10
300/300 [=====] - 0s 996us/step - loss: 0.0012
Epoch 8/10
300/300 [=====] - 1s 2ms/step - loss: 9.4196e-04
Epoch 9/10

```

```

300/300 [=====] - 0s 1ms/step - loss: 7.6574e-04
Epoch 10/10
300/300 [=====] - 0s 845us/step - loss: 6.3096e-04
300/300 [=====] - 0s 356us/step
100/100 [=====] - 0s 376us/step
Average Training Score: 0.024622411715387073
Average Validation Score: 0.030026559600401764

# Test set
y_test_pred = model.predict(X_test)
test_score = np.sqrt(mean_squared_error(y_test, y_test_pred))
print(f"Test Score: {test_score}")

100/100 [=====] - 0s 353us/step
Test Score: 0.02799366725417523

```